Contents of Lecture 3

- Threads in Java

- Synchronized methods and blocks

- Lock and Trylock in Java

- `volatile` in Java vs in C

- Pthreads

Creating Java Threads

- Either extend the `Thread` class or implement the `Runnable` interface.

- Your thread needs a `public void run()` method.

- Don't call `run` — you should instead call `start`.

- If `main` calls `run`, `main` has to do the work itself...

- To wait for a thread to terminate, you use `join` in a try-catch block.

The Work class

```java
class Work extends Thread {

        Node                    excess; // nodes with excess preflow
        Node                    s;      // source node
        Node                    t;      // sink node

        public void run()
        {
                while (excess != null) {
                }
        }
}
```

An incomplete preflow function

```
int preflow(int s, int t, int nthread)
{
        work                  = new Work[nthread];
        for (i = 0; i < nthread; ++i)
                work[i] = new Work(node[s], node[t]);

        for (i = 0; i < nthread; ++i)
                work[i].start();

        for (i = 0; i < nthread; ++i) {
                try {
                        work[i].join();
                } catch (Exception e) {
                        System.out.println("" + e);
                }
        }
}
```

Java synchronization

- A program which uses synchronization properly to avoid data races is said to be **thread safe**.

- Every Java object has a lock. Before entering a method declared `synchronized`, the JVM checks if the calling thread is the owner of the lock.

- If no thread owned the lock the calling thread becomes the owner at once.

- If the lock is owned by another thread, the calling thread is blocked and is put into an **entry set** for the object.

- When the lock is released, some thread in the entry set is resumed and becomes the new lock owner.

Synchronized blocks

- Not only methods can be synchronized.

```
void push(Node u, Node v, Edge e)
{
        synchronized (u) {
        }
}
```

- The behaviour is the same, that the object's lock is attempted to be taken

Locking a useless object

- Suppose a thread succeeds in taking an object but finds it useless
- For instance a queue which is empty
- We would then like to wait for another thread to "fix" the object — such as putting something in the queue
- As it is, we have the lock and nobody can fix it
- We want to unlock and wait for the fix
- And be told when when we should check it out again
- A Java object has a condition variable for this

The Java Object Wait-Set

- There are two methods `wait()` and `notify()` for this
- In addition to the entry set there is also a **wait set**.
- Calling `wait()` releases the lock, blocks the thread, and puts it into the wait set.
- The purpose is to let some other thread fix the object so it becomes usable for the thread which calls `wait()`.
- Calling `notify()` wakes up one thread in the wait set if there was any there, puts it it into the entry set and sets its thread state to Runnable.
- Calling `notify()` does **not** release the lock.
- One can also call `notifyAll()` which wakes up all threads in the wait set.

More Details

- The Java object lock is a so called **recursive lock** which means that we can call other synchronized methods for the same object without being blocked by ourselves.

- A thread may own multiple object locks (i.e., call synchronized methods for different objects).

- A notification for an object with an empty wait set has no effect.

Using `wait()`

- A thread waiting in the wait set can be interrupted, ie another thread can call its `interrupt()` method. This results in the `InterruptedException` being thrown:

```
try {

        wait();

}
catch (InterruptedException ie) { /* ignore */ }
```

ReentrantLock

- Sometimes it is more convenient to use an explicit lock and not the "implicit" / "internal" lock that is used with `synchronized`

```
import java.util.concurrent.locks.ReentrantLock;

ReentrantLock                   x = new ReentrantLock();

x.lock();
x.unlock();
if (x.tryLock()) {
        x.unlock();
}
```

- If we cannot see a simple lock-order rule we can sometimes use tryLock
- Say you want to take some locks in a random order. Try one at a time and restart with the first if any fails.
- Can you see a simple rule to avoid deadlocks and tryLock for Lab 2?

volatile in Java

- The definition of `volatile` has changed since it was introduced.

- Initially, the accessing of a particular volatile attribute of an object was serialized, i.e., all threads saw these accesses in the same order.

- These accesses were, however, unrelated to accesses of other variables.

- If you updated some variables and then set a volatile flag to indicate you were done, your program was buggy since the accesses were not ordered, i.e., the updates may be seen after the new value of the flag!

volatile in Java

- Now `volatile` is more similar to `synchronized`

- When entering a synchronized block, or reading a `volatile` attribute, everything in the cache is conceptually made invalid.

- It is not in reality but think of it as all variables must be fetched from memory with their most recent values that have previously been written to memory.

- When leaving a synchronized block or writing a `volatile` attribute everything in the cache is, again conceptually, written to memory.

- The same concept applies to locks in Java (and C/C++):
  - Taking a lock fetches every new variable from memory
  - Unlocking it writes everything back to memory

- In addition, unlocking must wait for OK from memory that every thread knows about the writes

- Note that this is only conceptual otherwise extremely slow

volatile in C

- `volatile` in C tells the compiler that accesses to this variable should never be optimized in any way

- It is not at all used for multithreading

- Normal variables can be allocated a register in the CPU but never volatile variables

Introducing Pthreads

- POSIX stands for Portable Operating System Interface and is an API for UNIX programmers

- Pthreads, or POSIX Threads, is available on all UNIX machines, including Linux and MacOS X

- Pthreads are quite similar to Java threads

- They are enabled by:
  - `#include <pthread.h>` in the C source file
  - compiling with: `gcc file.c -pthread`
  - or: `clang file.c -pthread`

Getting started

- `#include <pthread.h>`

- `pthread_t` is the type of a thread

- Create threads using `pthread_create()`

- Wait for a thread using `pthread_join()`

- Terminate a thread using `pthread_exit()`

pthread_create() 1(2)

```
int pthread_create(
        pthread_t*              thread,         // output.
        const pthread_attr_t*   attr,           // input.
        void*                   (*work)(void*), // input.
        void*                   arg);           // input.


pthread_t                   thread;
int                         status;
struct { int a, b, c }   arg = { 1, 2, 3 };

status = pthread_create(&thread, NULL, work, &arg);
```

- The thread identifier is filled in by the call and attributes are optional
- The created thread runs the `work` function and then terminates
- A class is called a `struct` in C — no methods and everything public
- Typically multiple arguments are passed in a struct as above

pthread_create() 2(2)

```
int pthread_create(
        pthread_t*              thread,         // output.
        const pthread_attr_t*   attr,           // input.
        void*                   (*work)(void*), // input.
        void*                   arg);           // input.
```

- A zero return value from `pthread_create()` indicates success, and a nonzero describes an error printable with `perror`.

- `void*` is a void pointer and is similar to Java's `Object`

- So the `work` function can return any data.

- Calling `pthread_join` waits for the termination of another thread and also gives access to the returned void pointer.

pthread_join()

```
int pthread_join(
        pthread_t        thread,          // input.
        void**           result);         // output.
```

- The call causes the caller to wait for the termination of a thread.
- If non-NULL, the terminated thread's return value is stored in result.
- A thread can only be joined by one thread.
- In Lab 2 you don't need to return any value from a thread and can just use:

```
pthread_t        thread[t];
for (i = 0; i < t; i += 1)
    if (pthread_create(&thread[i], NULL, work, arg) != 0)
        error("pthread_create failed");

for (i = 0; i < t; i += 1)
    if (pthread_join(thread[i], NULL) != 0)
        error("pthread_join failed");
```

pthread_exit()

`void pthread_exit(void*);        // return value from work.`

- Either use this or a return from the work function to terminate a thread.

- At termination of the main thread using `exit` or return, all other threads are killed.

- After a thread has terminated, the Pthreads system waits until some other thread joins with it. Then the terminated thread's resources are recycled.

- If a thread will never be joined, it should have been detached so that the system can recycle resources.

pthread_detach()

```
void pthread_detach(pthread_t thread);  // recycle at exit.
```

- A thread can be detached from the beginning by specifying an attribute saying so at `pthread_create`.
- Or, any thread can call `pthread_detach`.
- A detached thread cannot be joined.

Terminating a Thread

- There are three ways to terminate a thread:

1. Return from the work function.

2. The thread can call the function
   `void pthread_exit(void* value)`. The Standard C Library
   function `void exit(int status)` should normally not be used since
   it terminates the entire program.

3. Calling the function `int pthread_cancel(pthread_t thread)`
   makes a request to terminate the specified thread. See cancellation
   below.

- The first two are used by the thread itself and the third is used to stop
  another thread.

- Stopping another thread can be useful e.g. when a user has hit a
  "Cancel" button or another thread already has found a winning chess
  move.

Thread Cancellation Overview

- Simply terminating a thread can be disastrous if for example it has locked a mutex and is modifying shared data.

- Therefore the `pthread_cancel` simply requests that the thread should terminate. To actually know that the thread has terminated, it must be joined with.

- A thread that received a cancellation request is informed about this fact at certain points in the program, called **cancellation points**.

- The termination of the thread is started when it comes to such a cancellation point, if it has a **pending** cancellation request.

- A thread can install a function that is executed before the thread actually terminates.

- It is possible to allow cancellation at any time — see below.

Cancellation State and Type

- For cancellation, a thread has two variables, each with two possible values.

- The variables cannot be accessed directly but only through function calls.

- They are:
  - **State** — cancellation is either enabled or disabled.
  - **Type** — cancellation is either asynchronous or deferred.

- These result in three different cancellation modes:
  1. **Disabled** — any cancellation request received is saved until cancellation is enabled in the future.
  2. **Deferred** — cancellation is started at a cancellation point if there is a pending cancellation request.
  3. **Asynchronous** — cancellation can start at any time.

Modifying the Cancellation Mode

- The functions to modify the cancellation mode returns the thread's old value of the respective variable.

- `int pthread_setcancelstate(int state, int* old);`
  The state must be one of:
  - `PTHREAD_CANCEL_ENABLE`
  - `PTHREAD_CANCEL_DISABLE`

- `int pthread_setcanceltype(int type, int* old);` The type must be one of:
  - `PTHREAD_CANCEL_DEFERRED`
  - `PTHREAD_CANCEL_ASYNCHRONOUS`

- The default mode for new threads is **deferred**.

Cancellation Points

- A number of functions are cancellation points, including

`pthread_cond_wait`	`pthread_cond_timedwait`
`pthread_testcancel`	`pthread_join`
`close`	`creat`
`open`	`read`
`system`	`wait`
`waitpid`	`write`

- POSIX guarantees that the above (and some others) are cancellation points.
- Another list contains possible cancellation points, including

`printf`	`scanf`
`fopen`	`fclose`

- ISO C and POSIX functions not on any of those lists are guaranteed not to be cancellation points.

Receiving a Cancellation Request

- Thus if a cancellation request is received while cancellation is disabled, the request is simply blocked until it is enabled again.

- A pending request is delivered when the thread comes to a function which is a cancellation point.

- Changing the cancellation mode is **not** a cancellation point.

- At a cancellation point the thread first executes any installed cleanup handler (see below) and then terminates the thread.

- The return value from a cancelled thread is `PTHREAD_CANCELED` which thus is a valid value of a void pointer.

Installing Cleanup Handlers

- Each thread has a stack of cleanup handlers.

- They are installed with the function:
  `void pthread_cleanup_push(void (*func)(void*), void* arg);`

- The argument `arg` will be passed to `func` when it is executed.

- To remove a cleanup handler, use the function:
  `void pthread_cleanup_pop(int execute);`

- If the argument `execute` is nonzero, the cleanup handler will first be executed and then popped.

- When a thread is about to be terminated, all cleanup handlers on the stack are executed, starting with what is on the top of the stack.

Execution of Cleanup Handlers

- There are three situations when one or all cleanup handlers are executed:

  1. When a thread is being terminated due to a cancellation.
  2. When a thread is being terminated due to it has called `pthread_exit`.
  3. When it has called `pthread_cleanup_pop` with a nonzero parameter.

- In the last case, only one cleanup handler is executed, as we just saw.

Synchronization in Pthreads

- Pthreads has three main primitives for synchronization:
  - mutex
  - condition variable
  - barrier

Avoid synchronization!

- Ideally a parallel program needs no synchronization.

- Synchronization and therefore data communication between threads/caches take time.

- Some problems can be divided into suitable tasks statically.

- However, a common problem if T tasks are statically assigned to P threads is that some tasks take more time and therefore there becomes an inbalance in the work load, i.e. some threads take much longer time than the others.

Mutex

- A POSIX Threads **mutex** is a lock with a sleep queue

- The type is `pthread_mutex_t` and the most important functions related to it are:

- `pthread_mutex_init`

- `pthread_mutex_destroy`

- `pthread_mutex_lock`

- `pthread_mutex_trylock`

- `pthread_mutex_unlock`

- All five take a pointer to a `pthread_mutex_t`, and `pthread_mutex_init` also takes a pointer to attributes, which may be `NULL`.

Lock a mutex twice

- Trying to lock the same mutex multiple times does not work by default:

```
pthread_mutex_t A;

pthread_mutex_init(&A, NULL);
pthread_mutex_lock(&A);
pthread_mutex_lock(&A);
```

Recursive Mutex

- We must initialize the mutex as follows for this:

```
pthread_mutex_t A;
pthread_mutexattr_t      attr;

pthread_mutexattr_init(&attr);
pthread_mutexattr_settype(&attr, PTHREAD_MUTEX_RECURSIVE);
pthread_mutex_init(&A, &attr);

pthread_mutex_lock(&A);
pthread_mutex_lock(&A);
```

Worklists

- Assume we have a number of tasks to be processed.

- We put the tasks in lists and create threads which take tasks from the lists and process them.

- Concurrently adding or removing of items in the lists means the lists must be protected.

- Two alternatives:
  1. Put a mutex lock in each list head, i.e. protect the data.
  2. Use a common mutex for all lists, i.e. protect the code.

- Which is best depends on the application. There can be more concurrency if each list head has its own lock, at the cost of memory...

Condition variables in Pthreads

- As in Java, a condition variable lets a thread wait for something to happen in the future, and another thread to inform it that it has happened.

- For example: a worker thread can wait for a task being inserted in the list and another thread can signal any waiting thread that it just has inserted a new task.

- The condition variable type is: `pthread_cond_t`.

- In addition to initialization and destruction functions the main functions are:
  - `pthread_cond_wait` — causes calling thread to wait
  - `pthread_cond_signal` — wakes up one waiting thread
  - `pthread_cond_broadcast` — wakes up all waiting threads

Pthread mutex and condition variable

- Suppose you have locked a mutex and want to wait

- You need to both unlock and wait

- So that is done in one function atomically

```
pthread_cond_wait(&cond, &mutex);
```

- If they were two separate functions we would have problems

- Unlocking first could miss a signal

- (and waiting first could not unlock the mutex...)

More Details

- One mutex may be used for multiple condition variables

- Two threads wanting to wait for the same condition variable must use the same mutex.

- It is OK to signal a condition variable without having locked the corresponding mutex, but not so common.

Predicate, Condition Variable, and Mutex

- The logic expression in the C code which decides whether a thread should wait on the condition variable is called the **predicate** associated with the condition variable.

- The predicate is computed from shared data which different threads can modify, and therefore that data must be protected using a mutex.

- For example, the predicate may be computed by a boolean function:

```
bool empty(buffer_t* buffer);
```

- The predicate should always be tested in a loop and not in an if-statement.

Why You Need A Loop

- You should write your code like this.

```
pthread_mutex_lock(&mutex);
while (!predicate())
        pthread_cond_wait(&cond, &mutex);
/* do something... */
pthread_mutex_unlock(&mutex);
```

- There are at least three reasons for doing so:
  1. **Intercepted wakeups:** Another thread might have locked the mutex before.
  2. **Loose predicates:** This is a kind of predicate that says "the predicate may be true (but check before relying on it)."
  3. **Spurious wakeups:** This is very uncommon and is essentially an error that a thread was woke up without any good reason.

Intercepted Wakeups

- Should you signal a condition variable before or after you unlock its associated mutex (in case you have it locked) ???

- If you signal first, then the woke up thread will immediately try to lock the mutex and find it locked and wait again, now instead on a mutex, causing unnecessary synchronization overhead, both in the form of instructions and cache misses

- There may already be another thread waiting for the mutex which then later gets the mutex first

- If you unlock first, there may be a higher chance another thread takes the lock before the thread you wake up.

- It means the predicate might not longer be true after the other thread has unlocked the lock and it is your turn.

- This is called an intercepted wakeup.

- Due to intercepted wakeups, you must check the condition in a loop.

Loose Predicates

- It may be more convenient and/or efficient to say "you might have something interesting to check out" rather promising something.

- The woke up thread then must itself determine if there really was something for it, or whether it should continue waiting.

- This is called a loose predicate.

Spurious Wakeups

- When you hit CTRL-C to terminate a program you send a so called UNIX signal to it. This use of the word signal has nothing to do with the signal function of condition variables.

- When a thread receives a UNIX signal and was in the UNIX kernel waiting for a system call to complete, that system call is terminated and returns with the error code `EINTR`.

- Some UNIX signals are sent to all threads of a running program (called a **process**) and a thread waiting on a condition variable, i.e. in a system call on UNIX will thus be interrupted to handle the signal.

- An interrupted system call is not resumed but the application proceeds after it has returned.

- In principle a system call used by `pthread_cond_wait` could be interrupted and result in a spurious wakeup, but that is not the behaviour on Linux which uses the **futex** system call described below.

Implementation of Linux Native Pthreads Library

- Should a mutex lock involve the Linux kernel?

- Preferably not because it takes a lot of time

- On Linux there is a low-level synchronization primitive called **futex** which is used to implement the Pthreads library.

- Two good C libraries: gnu libc (glibc) and musl

- Linux command to show system calls: `strace -f ./a.out`

- `-f` tells it to trace threads and other programs it has *forked* (created)

User Level Locking with Futex in Linux

- Originates from IBM Research and the IBM Linux Technology Center.

- Implemented in the GNU C Library and in the Linux kernel, since version 2.5.7.

- The lock variable is in user space in shared memory and there is a corresponding wait queue for a lock in the kernel.

- The fast case is when there is no contention for the lock and therefore the kernel needs not be involved.

- The lock is manipulated in user space with atomic instructions (or the equivalent).

Initialization of a Pthread Mutex

```
pthread_mutex_t          mutex = PTHREAD_MUTEX_INITIALIZER;
```

- This initializes the mutex with default attributes.

- `PTHREAD_MUTEX_INITIALIZER` is a constant expression, meaning we can initialize a mutex like this at file scope (static storage, ie a `static` or global variable).

- If allocated by eg malloc, then `pthread_mutex_init()` should be called.

- After usage, `pthread_mutex_destroy` should be called for a mutex, and then its memory should be deallocated using `free`, if appropriate.

pthread_cond_timedwait

- To wait on a condition variable with a time out, use `pthread_cond_timedwait`.

  ```
  int pthread_cond_timedwait(
          pthread_cond_t*,
          pthread_mutex_t*,
          struct time_spec*);
  ```

- The time is absolute time and to wait eg for at most 3 seconds, one can use:

  ```
  timeout.tv_sec = time(NULL) + 3;
  timeout.tv_nsec = 0;
  ```

- If there is a time out, the return value is `ETIMEDOUT`.

Pthreads barriers

- A barrier is used to let all threads work in a more synchronous way.

- All threads must reach `pthread_barrier_wait` before any can proceed beyond it.

```
int pthread_barrier_init(
        pthread_barrier_t*     barrier,
        pthread_barrierattr_t*  attr,
        unsigned int            count);


int pthread_barrier_destroy(pthread_barrier_t* barrier);


int pthread_barrier_wait(pthread_barrier_t* barrier);
```

Initialization in sequential programs

- The usual sequential way to initialize is to do like this:

```c
#include <stdbool.h>
void f(void)
{
        static bool initialized = false;
        if (!initialized) {
                init();
                initialized = true;
        }
        ...
}
```

- Might not work in a multithreaded program!

Initialization in multithreadthreaded programs

```c
#include <stdbool.h>
pthread_mutex_t  init_lock;
void f(void)
{
        static bool initialized = false;
        pthread_mutex_lock(&init_lock);
        if (!initialized) {
                init();
                initialized = true;
        }
        pthread_mutex_unlock(&init_lock);
}
```

pthread_once

- One can write `pthread_once_t once = PTHREAD_ONCE_INIT;`.

- The `once` variable can have static storage duration.

- The function

  ```
  int pthread_once(pthread_once_t*, void (*)(void));
  ```

  is used to call a function once. It takes a pointer to a "once" variable and a function to execute for the initialization.

- If another thread executes the same call after the first is done, nothing will happen.

- If it instead calls it during the first call, the second thread will wait until the first call is done.

Thread attributes

- Examples of attributes which can be set:
  - Whether another thread can join with a particular thread (portable).
  - Stack address
  - Stack size (not portable)

- A thread which is not joinable is "detached" which means the resource used by the thread are recycled immediately when the thread terminates.

- The joinable attribute can be set to one of
  - `PTHREAD_CREATE_JOINABLE`, or
  - `PTHREAD_CREATE_DETACHED`.

- An initially joinable thread can make itself detached but not vice versa.